

K-Nearest Neighbors

K Sri Rama Murty

IIT Hyderabad

`ksrm@ee.iith.ac.in`

K-Nearest Neighbor Classifier

K-Nearest Neighbor Classifier

- Classification is a supervised learning problem

K-Nearest Neighbor Classifier

- Classification is a supervised learning problem
- In classification, the target variable admits a discrete set of values.
eg. $[(x_1=6 \text{ ft}, x_2=65 \text{ kg}), y=A], [(x_1=3.4 \text{ ft}, x_2=18 \text{ kg}), y=K]$

K-Nearest Neighbor Classifier

- Classification is a supervised learning problem
- In classification, the target variable admits a discrete set of values.
eg. $[(x_1=6 \text{ ft}, x_2=65 \text{ kg}), y=A], [(x_1=3.4 \text{ ft}, x_2=18 \text{ kg}), y=K]$
- Given a set of labeled data $[\mathbf{x}_n, y_n], n = 1, 2, \dots, N$ for training, classify a new instance \mathbf{x}_0

K-Nearest Neighbor Classifier

- Classification is a supervised learning problem
- In classification, the target variable admits a discrete set of values.
eg. $[(x_1=6 \text{ ft}, x_2=65 \text{ kg}), y=A], [(x_1=3.4 \text{ ft}, x_2=18 \text{ kg}), y=K]$
- Given a set of labeled data $[\mathbf{x}_n, y_n], n = 1, 2, \dots, N$ for training, classify a new instance \mathbf{x}_0
- KNN classification algorithm

K-Nearest Neighbor Classifier

- Classification is a supervised learning problem
- In classification, the target variable admits a discrete set of values.
eg. $[(x_1=6 \text{ ft}, x_2=65 \text{ kg}), y=A], [(x_1=3.4 \text{ ft}, x_2=18 \text{ kg}), y=K]$
- Given a set of labeled data $[\mathbf{x}_n, y_n]$, $n = 1, 2, \dots, N$ for training, classify a new instance \mathbf{x}_0
- KNN classification algorithm
 - Evaluate distance of \mathbf{x}_0 to all points \mathbf{x}_n , $n = 1, 2, \dots, N$ in training set

$$d_n = \|\mathbf{x}_n - \mathbf{x}_0\|_p \qquad d_n = 1 - \frac{\mathbf{x}_n^T \mathbf{x}_0}{\|\mathbf{x}_n\| \|\mathbf{x}_0\|}$$

K-Nearest Neighbor Classifier

- Classification is a supervised learning problem
- In classification, the target variable admits a discrete set of values.
eg. $[(x_1=6 \text{ ft}, x_2=65 \text{ kg}), y=A], [(x_1=3.4 \text{ ft}, x_2=18 \text{ kg}), y=K]$
- Given a set of labeled data $[\mathbf{x}_n, y_n], n = 1, 2, \dots, N$ for training, classify a new instance \mathbf{x}_0
- KNN classification algorithm
 - Evaluate distance of \mathbf{x}_0 to all points $\mathbf{x}_n, n = 1, 2, \dots, N$ in training set

$$d_n = \|\mathbf{x}_n - \mathbf{x}_0\|_p \qquad d_n = 1 - \frac{\mathbf{x}_n^T \mathbf{x}_0}{\|\mathbf{x}_n\| \|\mathbf{x}_0\|}$$

- Sort the distances and identify K nearest neighbors - X_{NN}




K-Nearest Neighbor Classifier

- Classification is a supervised learning problem
- In classification, the target variable admits a discrete set of values.
eg. $[(x_1=6 \text{ ft}, x_2=65 \text{ kg}), y=A], [(x_1=3.4 \text{ ft}, x_2=18 \text{ kg}), y=K]$
- Given a set of labeled data $[\mathbf{x}_n, y_n], n = 1, 2, \dots, N$ for training, classify a new instance \mathbf{x}_0
- KNN classification algorithm
 - Evaluate distance of \mathbf{x}_0 to all points $\mathbf{x}_n, n = 1, 2, \dots, N$ in training set

$$d_n = \|\mathbf{x}_n - \mathbf{x}_0\|_p \qquad d_n = 1 - \frac{\mathbf{x}_n^T \mathbf{x}_0}{\|\mathbf{x}_n\| \|\mathbf{x}_0\|}$$



- Sort the distances and identify K nearest neighbors - X_{NN}
- Assign output label y_0 for test point \mathbf{x}_0 from labels of y_{NN} - Voting

Geometry of KNN Decision Boundaries ($k=1$)




Steel-plate classification

Geometry of KNN Decision Boundaries (k=1)




Steel-plate classification

Geometry of KNN Decision Boundaries (k=1)

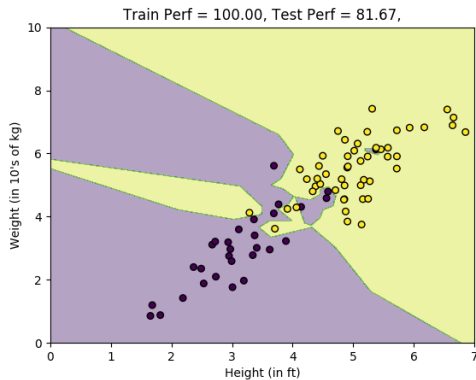


Steel-plate classification




Decision boundaries are perpendicular bisectors of the lines joining nearest competitors!

KNN Classifier (Kid vs Adult) $k = 1$

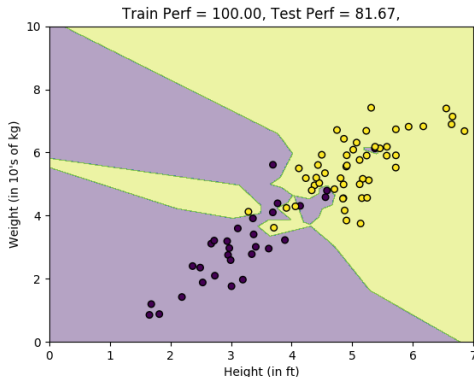


KNN Classifier (Kid vs Adult) $k = 1$




- 50 points from adult-class and 30 points from kid-class

KNN Classifier (Kid vs Adult) $k = 1$




- 50 points from adult-class and 30 points from kid-class
- Big difference between train and test performances ($k = 1$ is a overfit)

KNN Classifier (Kid vs Adult) $k = 1$





- 50 points from adult-class and 30 points from kid-class
- Big difference between train and test performances ($k = 1$ is a overfit)
- In a KNN-classifier, the decision regions are typically nonconvex




Effect of K on Performance







Effect of K on Performance




Effect of K on Performance





Effect of K on Performance



Why not Weight in kg





Why not Weight in kg





Why not Weight in kg

- Distance between two points \mathbf{x} and \mathbf{y} is

$$d_{xy} = \sqrt{(x_h - y_h)^2 + (x_w - y_w)^2}$$



Why not Weight in kg





- Distance between two points \mathbf{x} and \mathbf{y} is

$$d_{xy} = \sqrt{(x_h - y_h)^2 + (x_w - y_w)^2}$$

- Contribution of a dimension to distance is proportional to its variance

Why not Weight in kg





- Distance between two points \mathbf{x} and \mathbf{y} is

$$d_{xy} = \sqrt{(x_h - y_h)^2 + (x_w - y_w)^2}$$

- Contribution of a dimension to distance is proportional to its variance
- When weight is in kg, distance is dominated by it (horizontal boundary)

Why not Weight in kg



- Distance between two points \mathbf{x} and \mathbf{y} is



$$d_{xy} = \sqrt{(x_h - y_h)^2 + (x_w - y_w)^2}$$

- Contribution of a dimension to distance is proportional to its variance
- When weight is in kg, distance is dominated by it (horizontal boundary)
- Normalize each dim. to unit variance

$$\tilde{x}_w = \frac{x_w - \mu_w}{\sigma_w}$$

$$\sigma_w^2 = \frac{1}{N} \sum_{n=1}^N (x_{nw} - \mu_w)^2 \quad \mu_w = \frac{1}{N} \sum_{n=1}^N x_{nw}$$

Why not Weight in kg



- Distance between two points \mathbf{x} and \mathbf{y} is
$$d_{xy} = \sqrt{(x_h - y_h)^2 + (x_w - y_w)^2}$$
- Contribution of a dimension to distance is proportional to its variance
- When weight is in kg, distance is dominated by it (horizontal boundary)
- Normalize each dim. to unit variance

$$\tilde{x}_w = \frac{x_w - \mu_w}{\sigma_w}$$

$$\sigma_w^2 = \frac{1}{N} \sum_{n=1}^N (x_{nw} - \mu_w)^2 \quad \mu_w = \frac{1}{N} \sum_{n=1}^N x_{nw}$$

- This process is called *data whitening*

k -Nearest Neighbor (k -NN) Regression

k -Nearest Neighbor (k -NN) Regression

- Regression is a supervised learning problem

k -Nearest Neighbor (k -NN) Regression

- Regression is a supervised learning problem
- In regression, the target variable admits continuous values

S.No	1	2	3	4	5	6	7	8	9	10	t
x	5.2	4.8	5.8	5.7	5.4	5.1	4.9	5.3	5.9	4.5	5.6
y	58	53	63	65	60	50	52	55	61	49	?

k -Nearest Neighbor (k -NN) Regression

- Regression is a supervised learning problem
- In regression, the target variable admits continuous values

S.No	1	2	3	4	5	6	7	8	9	10	t
x	5.2	4.8	5.8	5.7	5.4	5.1	4.9	5.3	5.9	4.5	5.6
y	58	53	63	65	60	50	52	55	61	49	?

- Average of the k -nearest neighbors in the training data.

k -Nearest Neighbor (k -NN) Regression

- Regression is a supervised learning problem
- In regression, the target variable admits continuous values

S.No	1	2	3	4	5	6	7	8	9	10	t
x	5.2	4.8	5.8	5.7	5.4	5.1	4.9	5.3	5.9	4.5	5.6
y	58	53	63	65	60	50	52	55	61	49	?

- Average of the k -nearest neighbors in the training data.
- ($k = 1$) NN - $\{(5.7, 65)\}$: $y_{11} = 65$ (overfit - high variance)

k -Nearest Neighbor (k -NN) Regression

- Regression is a supervised learning problem
- In regression, the target variable admits continuous values

S.No	1	2	3	4	5	6	7	8	9	10	t
x	5.2	4.8	5.8	5.7	5.4	5.1	4.9	5.3	5.9	4.5	5.6
y	58	53	63	65	60	50	52	55	61	49	?

- Average of the k -nearest neighbors in the training data.
- ($k = 1$) NN - $\{(5.7, 65)\}$: $y_{11} = 65$ (overfit - high variance)
- ($k = 3$) NN - $\{(5.7, 65), (5.8, 63), (5.4, 60)\}$: $y_{11} = 62.66$

k -Nearest Neighbor (k -NN) Regression

- Regression is a supervised learning problem
- In regression, the target variable admits continuous values

S.No	1	2	3	4	5	6	7	8	9	10	t
x	5.2	4.8	5.8	5.7	5.4	5.1	4.9	5.3	5.9	4.5	5.6
y	58	53	63	65	60	50	52	55	61	49	?

- Average of the k -nearest neighbors in the training data.
- ($k = 1$) NN - $\{(5.7, 65)\}$: $y_{11} = 65$ (overfit - high variance)
- ($k = 3$) NN - $\{(5.7, 65), (5.8, 63), (5.4, 60)\}$: $y_{11} = 62.66$
- ($k = 10$) assigns average weight irrespective of test height (high bias)

k -Nearest Neighbor (k -NN) Regression

- Regression is a supervised learning problem
- In regression, the target variable admits continuous values

S.No	1	2	3	4	5	6	7	8	9	10	t
x	5.2	4.8	5.8	5.7	5.4	5.1	4.9	5.3	5.9	4.5	5.6
y	58	53	63	65	60	50	52	55	61	49	?

- Average of the k -nearest neighbors in the training data.
- ($k = 1$) NN - $\{(5.7, 65)\}$: $y_{11} = 65$ (overfit - high variance)
- ($k = 3$) NN - $\{(5.7, 65), (5.8, 63), (5.4, 60)\}$: $y_{11} = 62.66$
- ($k = 10$) assigns average weight irrespective of test height (high bias)
- KNN-regression gives equal importance to all neighbors

k -Nearest Neighbor (k -NN) Regression

- Regression is a supervised learning problem
- In regression, the target variable admits continuous values

S.No	1	2	3	4	5	6	7	8	9	10	t
x	5.2	4.8	5.8	5.7	5.4	5.1	4.9	5.3	5.9	4.5	5.6
y	58	53	63	65	60	50	52	55	61	49	?

- Average of the k -nearest neighbors in the training data.
- ($k = 1$) NN - $\{(5.7, 65)\}$: $y_{11} = 65$ (overfit - high variance)
- ($k = 3$) NN - $\{(5.7, 65), (5.8, 63), (5.4, 60)\}$: $y_{11} = 62.66$
- ($k = 10$) assigns average weight irrespective of test height (high bias)
- KNN-regression gives equal importance to all neighbors
- Assign higher weightage to closer neighbors

Inverse Distance Weighted k -NN

Inverse Distance Weighted k -NN

- Weighted k -NN regression: the unknown quantity y_t is estimated as

$$y_t = \frac{\sum_{n=0}^N k(x_n, x_t) y_n}{\sum_{n=0}^N k(x_n, x_t)}$$

Inverse Distance Weighted k -NN

- Weighted k -NN regression: the unknown quantity y_t is estimated as

$$y_t = \frac{\sum_{n=0}^N k(x_n, x_t) y_n}{\sum_{n=0}^N k(x_n, x_t)}$$

- $k(x_n, x_t)$ denotes similarity between x_n and x_t

Inverse Distance Weighted k -NN

- Weighted k -NN regression: the unknown quantity y_t is estimated as

$$y_t = \frac{\sum_{n=0}^N k(x_n, x_t) y_n}{\sum_{n=0}^N k(x_n, x_t)}$$

- $k(x_n, x_t)$ denotes similarity between x_n and x_t
- Similarity (or inverse distance) as weight to give preference to NNs

Inverse Distance Weighted k -NN

- Weighted k -NN regression: the unknown quantity y_t is estimated as

$$y_t = \frac{\sum_{n=0}^N k(x_n, x_t) y_n}{\sum_{n=0}^N k(x_n, x_t)}$$

- $k(x_n, x_t)$ denotes similarity between x_n and x_t
- Similarity (or inverse distance) as weight to give preference to NNs
- ($k = 3$) NN - $\{(10, 65), (5, 63), (5, 60)\}$: $y_{11} = 63.25$

Inverse Distance Weighted k -NN

- Weighted k -NN regression: the unknown quantity y_t is estimated as

$$y_t = \frac{\sum_{n=0}^N k(x_n, x_t) y_n}{\sum_{n=0}^N k(x_n, x_t)}$$

- $k(x_n, x_t)$ denotes similarity between x_n and x_t
- Similarity (or inverse distance) as weight to give preference to NNs
- $(k = 3)$ NN - $\{(10, 65), (5, 63), (5, 60)\}$: $y_{11} = 63.25$
- Number of neighbors k is not very critical.

Inverse Distance Weighted k -NN




- Weighted k -NN regression: the unknown quantity y_t is estimated as

$$y_t = \frac{\sum_{n=0}^N k(x_n, x_t) y_n}{\sum_{n=0}^N k(x_n, x_t)}$$




- $k(x_n, x_t)$ denotes similarity between x_n and x_t
- Similarity (or inverse distance) as weight to give preference to NNs
- $(k = 3)$ NN - $\{(10, 65), (5, 63), (5, 60)\}$: $y_{11} = 63.25$
- Number of neighbors k is not very critical.
- Similarity measure $k(x_n, x_t)$ is an important design choice

KNN-Regression Illustration


Uniform weights for the neighbors




Inverse distance weights for the neighbors



Robotic Arm Control - ME Approach



Robotic Arm Control - ME Approach




Solving for the forces:

$$F_{\theta_1} = ((M_1 + M_2)L_1^2 + M_2L_2^2 + 2M_2L_1L_2\cos\theta_2)\ddot{\theta}_1 + (M_2L_2^2 + M_2L_1L_2\cos\theta_2)\ddot{\theta}_2 - M_2L_1L_2\sin\theta_2(2\dot{\theta}_1\dot{\theta}_2 + \dot{\theta}_2^2) - (M_1 + M_2)gL_1\sin\theta_1 - M_2gL_2\sin(\theta_1 + \theta_2)$$

$$F_{\theta_2} = (M_2L_2^2 + M_2L_1L_2\cos\theta_2)\ddot{\theta}_1 + M_2L_2^2\ddot{\theta}_2 - M_2L_1L_2\sin(\theta_2)\dot{\theta}_1\dot{\theta}_2 - M_2gL_2\sin(\theta_1 + \theta_2)$$

Robotic Arm Control - ME Approach



Solving for the forces:

$$F_{\theta_1} = ((M_1 + M_2)L_1^2 + M_2L_2^2 + 2M_2L_1L_2\cos\theta_2)\ddot{\theta}_1 + (M_2L_2^2 + M_2L_1L_2\cos\theta_2)\ddot{\theta}_2 - M_2L_1L_2\sin\theta_2(2\dot{\theta}_1\dot{\theta}_2 + \dot{\theta}_2^2) - (M_1 + M_2)gL_1\sin\theta_1 - M_2gL_2\sin(\theta_1 + \theta_2)$$

$$F_{\theta_2} = (M_2L_2^2 + M_2L_1L_2\cos\theta_2)\ddot{\theta}_1 + M_2L_2^2\ddot{\theta}_2 - M_2L_1L_2\sin(\theta_2)\dot{\theta}_1\dot{\theta}_2 - M_2gL_2\sin(\theta_1 + \theta_2)$$

- Force applied on robotic arm needs to be controlled
- Error between current and desired positions determine the forces
- F_{θ_1} and F_{θ_2} can be computing by solving Lagrangian using kinetic and potential energies
- Too complicated, and may not work
- However, we can conclude that

$$F = g(\theta_1, \theta_2, \dot{\theta}_1, \dot{\theta}_2, \ddot{\theta}_1, \ddot{\theta}_2)$$

Robotic Arm Control - KNN Approach

θ_1	θ_2	$\dot{\theta}_1$	$\dot{\theta}_2$	$\ddot{\theta}_1$	$\ddot{\theta}_2$	F_{θ_1}	F_{θ_2}

- From the current and desired positions, estimate the angles and their derivatives.
- Look for the $k(=1)$ -nearest match in the table, and pick forces.
- If succeeded in pitching, append the current configuration to the table
- As the table size increases, the robot skill improves!


Robot Training - Testing

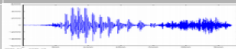
Speech Synthesis (Unit Selection)

- Regress speech waveform from the sequence of words
- Create a table of phonemes and their waveform from training data
- During synthesis, retrieve the waveforms for the sequence of phones, and concatenate them

Speech Synthesis (Unit Selection)

- Regress speech waveform from the sequence of words
- Create a table of phonemes and their waveform from training data
- During synthesis, retrieve the waveforms for the sequence of phones, and concatenate them



Phone	Prev. Phone	Next Phone	Phone position	Word Position	Duration	Waveform
/sh/	/a/	/i/	3/6	5/7	65 ms	

2h Data - 13h Data

KNN Summary

- **Pros**

- Simple to implement
- Flexible to feature/distance function choices (no differentiability)
- Naturally handles multiclass data
- Does well even on complex tasks with enough data

KNN Summary

- **Pros**

- Simple to implement
- Flexible to feature/distance function choices (no differentiability)
- Naturally handles multiclass data
- Does well even on complex tasks with enough data

- **Cons**

- Needs to access entire training data for every inference
- Computational complexity grows linearly with training data $O(ND)$

KNN Summary

• Pros

- Simple to implement
- Flexible to feature/distance function choices (no differentiability)
- Naturally handles multiclass data
- Does well even on complex tasks with enough data

• Cons

- Needs to access entire training data for every inference
- Computational complexity grows linearly with training data $O(ND)$
- Distance function has to be chosen carefully
- The choice of k is empirical (no theoretical guidelines)

KNN Summary

• Pros

- Simple to implement
- Flexible to feature/distance function choices (no differentiability)
- Naturally handles multiclass data
- Does well even on complex tasks with enough data

• Cons

- Needs to access entire training data for every inference
- Computational complexity grows linearly with training data $O(ND)$
- Distance function has to be chosen carefully
- The choice of k is empirical (no theoretical guidelines)
- Nonconvex decision regions
- Cannot naturally ignore noninformative dimensions

KNN Summary

• Pros

- Simple to implement
- Flexible to feature/distance function choices (no differentiability)
- Naturally handles multiclass data
- Does well even on complex tasks with enough data

• Cons

- Needs to access entire training data for every inference
- Computational complexity grows linearly with training data $O(ND)$
- Distance function has to be chosen carefully
- The choice of k is empirical (no theoretical guidelines)
- Nonconvex decision regions
- Cannot naturally ignore noninformative dimensions
- Does not consider all the data at once, and hence, it cannot model the underlying process responsible for the observed data

Homework - Data Generation Models

- Let the distribution of human height be given by

$$X \sim \mathcal{N}(5.3, 1)$$

Let the relationship between height and weight be given by

$$Y = 10X + 5 + \epsilon$$

where $\epsilon \sim \mathcal{N}(0, 3)$. Evaluate the joint density function $p_{XY}(x, y)$, marginal density functions $p_X(x)$ & $p_Y(y)$, conditional density functions $p(x/y)$ & $p(y/x)$, and their corresponding expectations.

Homework

- Let the joint density function of human heights and weights follow a multivariate Gaussian given by

$$[X, Y] \sim \mathcal{N}(\boldsymbol{\mu}, \boldsymbol{\Sigma})$$

where $\boldsymbol{\mu} \in \mathbb{R}^2$ is the mean vector and $\boldsymbol{\Sigma} \in \mathbb{R}^{2 \times 2}$ is the covariance matrix. Evaluate the joint density function $p_{XY}(x, y)$, marginal density functions $p_X(x)$ & $p_Y(y)$, conditional density functions $p(x/y)$ & $p(y/x)$, and their corresponding expectations.

- Is it related to the earlier model?

Homework

- Let \mathbf{V} be a 2-D random vector drawn from a standard Gaussian, i.e.,

$$\mathbf{V} \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$$


Let height and weight be obtained by an affine transformation of \mathbf{V} as

$$\begin{bmatrix} X \\ Y \end{bmatrix} = \begin{bmatrix} a & b \\ c & d \end{bmatrix} \mathbf{V} + \begin{bmatrix} \mu_x \\ \mu_y \end{bmatrix}$$


Evaluate the joint density function $p_{XY}(x, y)$, marginal density functions $p_X(x)$ & $p_Y(y)$, conditional density functions $p(x/y)$ & $p(y/x)$, and their corresponding expectations.

- Is it related to two earlier models?

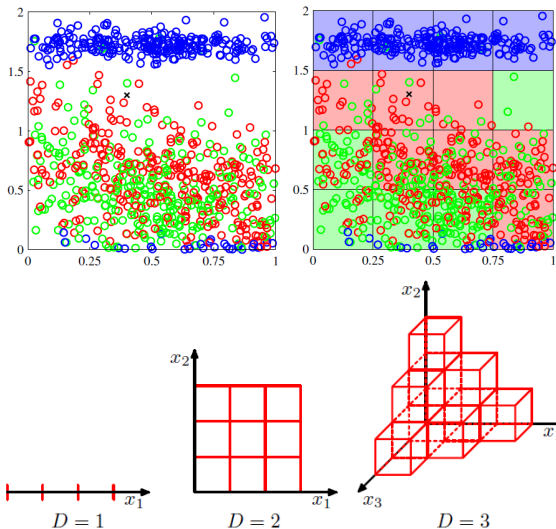
Curse of Dimensionality



Curse of Dimensionality




Curse of Dimensionality




Understanding High-Dimensional Space

2-D Gaussian scatter plot




Understanding High-Dimensional Space

2-D Gaussian scatter plot




3-D Gaussian scatter plot




Understanding High-Dimensional Space

2-D Gaussian scatter plot



3-D Gaussian scatter plot




1000-D Gaussian


?

Understanding High-Dimensional Space

2-D Gaussian scatter plot




3-D Gaussian scatter plot



1000-D Gaussian


?

Histogram of min and max of dimensions




Understanding High-Dimensional Space

2-D Gaussian scatter plot




3-D Gaussian scatter plot




1000-D Gaussian

?

Histogram of min and max of dimensions




Histogram of min and max of dimensions




Understanding High-Dimensional Space

2-D Gaussian scatter plot




3-D Gaussian scatter plot




1000-D Gaussian

?


Histogram of min and max of dimensions



Histogram of min and max of dimensions



Histogram of min and max of dimensions



Gaussian in High Dimensional Space


- The probability mass of a Gaussian distribution is concentrated in a thin shell

$$p(r)\Delta r = P\left[r - \frac{\Delta r}{2} < R < r + \frac{\Delta r}{2}\right]$$

Gaussian in High Dimensional Space

- The probability mass of a Gaussian distribution is concentrated in a thin shell

$$p(r)\Delta r = P\left[r - \frac{\Delta r}{2} < R < r + \frac{\Delta r}{2}\right]$$



Thank You !

Understanding High-Dimensional Space


- Most of the volume of a sphere is concentrated in thin shell near the surface.

$$\frac{V_D(1) - V_D(1 - \epsilon)}{V_D(1)} = 1 - (1 - \epsilon)^D$$

Understanding High-Dimensional Space

- Most of the volume of a sphere is concentrated in thin shell near the surface.

$$\frac{V_D(1) - V_D(1 - \epsilon)}{V_D(1)} = 1 - (1 - \epsilon)^D$$




Understanding High-Dimensional Space

- Most of the volume of a sphere is concentrated in thin shell near the surface.

$$\frac{V_D(1) - V_D(1 - \epsilon)}{V_D(1)} = 1 - (1 - \epsilon)^D$$

- The probability mass of a Gaussian distribution is concentrated in a thin shell

$$p(r)\Delta r = P\left[r - \frac{\Delta r}{2} < R < r + \frac{\Delta r}{2}\right]$$



Understanding High-Dimensional Space

- Most of the volume of a sphere is concentrated in thin shell near the surface.

$$\frac{V_D(1) - V_D(1 - \epsilon)}{V_D(1)} = 1 - (1 - \epsilon)^D$$

- The probability mass of a Gaussian distribution is concentrated in a thin shell

$$p(r)\Delta r = P\left[r - \frac{\Delta r}{2} < R < r + \frac{\Delta r}{2}\right]$$

