

Linear Models of Regression

K Sri Rama Murty

IIT Hyderabad

`ksrm@ee.iith.ac.in`

Regression

- Predict target variable(s) $t \in \mathbb{R}$ given D-dimensional input vector \mathbf{x}
- E.g. Weight estimation, Share market prediction, 3D image from 2D
- Target can be estimated as a linear combination of inputs

$$\hat{t} = y(\mathbf{x}, \mathbf{w}) = w_0 + w_1 x_1 + w_2 x_2 + \cdots w_D x_D = \mathbf{w}^T \mathbf{x}$$

$$\mathbf{x} = [1 \ x_1 \ x_2 \ \cdots \ x_D]^T \quad \mathbf{w} = [w_0 \ w_1 \ w_2 \ \cdots \ w_D]^T$$


- Determine the model parameters \mathbf{w} to minimize error on labeled training data

$$\mathcal{S} = \{(\mathbf{x}_1, t_1), (\mathbf{x}_2, t_2), \cdots (\mathbf{x}_N, t_N)\}$$


- Need to define a loss function for optimizing model parameters \mathbf{w}

Illustration of Linear Regression

Simple Linear Regression



Multiple Linear Regression
(2 Independent Variables (x_1, x_2))



Least Squares Criterion to Determine \mathbf{w}

- Estimated target for n^{th} sample in the dataset \mathcal{S}

$$\hat{t}_n = y(\mathbf{x}_n, \mathbf{w}) = \mathbf{w}^T \mathbf{x}_n \quad n = 1, 2, \dots, N$$

- Given the ground truth target t_n , the error in estimation

$$e_n = t_n - y_n \quad n = 1, 2, \dots, N$$

- For an arbitrary choice of parameters \mathbf{w} , overall error on training set

$$J(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N e_n^2$$

- Estimate \mathbf{w} to minimize loss function $J(\mathbf{w})$ on the training dataset \mathcal{S}

$$\mathbf{w}_* = \arg \min_{\mathbf{w}} J(\mathbf{w})$$

Estimating Optimal \mathbf{w}

- Let the input data be organized in the form of a $N \times (D + 1)$ matrix

$$\mathbf{X} = [\mathbf{x}_1 \ \mathbf{x}_2 \ \cdots \ \mathbf{x}_N]^T$$

- Estimated output vector \mathbf{y} for all data points is given by

$$\mathbf{y}_{N \times 1} = \mathbf{X}_{N \times D+1} \mathbf{w}_{D+1 \times 1}$$

- Loss function: Trace of the outer product of error vector $\mathbf{e} = \mathbf{t} - \mathbf{y}$


$$J(\mathbf{w}) = \frac{1}{2} \text{Tr}[\mathbf{e} \ \mathbf{e}^T]$$

- Equating derivative of loss function w.r.t \mathbf{w} to $\mathbf{0}$


$$\begin{aligned} \nabla_{\mathbf{w}} J(\mathbf{w}) &= \nabla_{\mathbf{w}} \mathbf{y} \ \nabla_{\mathbf{y}} \mathbf{e} \ \nabla_{\mathbf{e}} J(\mathbf{w}) \\ &= -\mathbf{X}^T (\mathbf{t} - \mathbf{X} \mathbf{w}) = \mathbf{0} \end{aligned}$$

$$\boxed{\mathbf{w} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{t}}$$

Geometric Interpretation of Least Squares



Error Surface of a Linear Neuron with Two Input Weights



- Given N examples, the target vector $\mathbf{t} \in \mathbb{R}^N$ and columns of $\mathbf{X} \in \mathbb{R}^N$
- Let \mathcal{S} denote a subspace spanned by columns of \mathbf{X} in N -dim space
- $\mathbf{y} = \mathbf{X}\mathbf{w} \in \mathcal{S}$, being a linear combination of columns of \mathbf{X}
- For the LS optimality criterion
 - \mathbf{y} is orthogonal projection of \mathbf{t} on \mathcal{S}
 - Error surface $J(\mathbf{w})$ is convex
 - Sim. to Wiener filter: $\mathbf{w} = \mathbf{R}_{xx}^{-1} \mathbf{r}_{xt}$
 - Also referred to as pseudo-inverse sol.

Homework

- Prove the following matrix derivatives



$$\nabla_{\mathbf{X}} \text{Tr}(\mathbf{XA}) = \mathbf{A}^T \quad \nabla_{\mathbf{X}} \text{Tr}(\mathbf{XA}) = \mathbf{A}^T \quad \nabla_{\mathbf{X}} \text{Tr}(\mathbf{AX}^T) = \mathbf{A}$$

$$\nabla_{\mathbf{X}} \text{Tr}(\mathbf{AXB}) = \mathbf{A}^T \mathbf{B}^T \quad \nabla_{\mathbf{X}} \text{Tr}(\mathbf{AX}^T \mathbf{B}) = \mathbf{BA}$$

$$\nabla_{\mathbf{X}} \text{Tr}(\mathbf{XAX}^T) = \mathbf{X}(\mathbf{A} + \mathbf{A}^T)$$

- True or False: Two nonzero, noncollinear vectors span \mathbb{R}^2 . If your answer is true express an arbitrary vector \mathbf{a} in terms of two noncollinear vectors \mathbf{v}_1 and \mathbf{v}_2 .
- Prove that there cannot be any other w_0 that gives a lower error than w_* to least squares formulation.
- Extend the linear regression concept to estimate multivariate target vectors $\mathbf{t}_n \in \mathbb{R}^K$

Nonlinear Input-Output Relations




- Polynomial curve fitting can be used to model nonlinear i/o relation

$$\begin{aligned}\hat{t} = y(x, \mathbf{w}) &= w_0 + w_1x + w_2x^2 + \cdots + w_Mx^M \\ &= \mathbf{w}^T \phi(x) \quad (\text{Model is linear in } \mathbf{w})\end{aligned}$$


- $\phi(.) : \mathbb{R}^1 \rightarrow \mathbb{R}^M$ - nonlinear transformation to higher dim. space

Kernel Examples


Polynomial



Gaussian



Logistic



$$\phi_j(x) = x^j$$

$$\phi_j(x) = \exp\left(-\frac{\|x - \mu_j\|^2}{2\sigma^2}\right)$$

$$\phi_j(x) = \sigma\left(\frac{x - \mu_j}{s}\right)$$

- Global vs Local kernels
 - Changes in one region of input space affect all other regions
 - Local kernels are preferable for functions with varying characteristics
- Explicit vs Implicit kernels
 - Explicit representation for $\phi(\mathbf{x})$ is available or not





Least Squares Regression in Kernel Space

- If t_n is nonlinearly related to \mathbf{x}_n , perform regression in kernel space.
- Let $\phi : \mathbb{R}^D \rightarrow \mathbb{R}^{M+1}$, $M > D$ is a nonlinear kernel mapping
- $\mathbf{x}_n = [x_{n1} \ x_{n2}]^T \in \mathbb{R}^2$ can be mapped using 2nd order polynomial kernel as $\phi(\mathbf{x}_n) = [1 \ x_{n1} \ x_{n2} \ x_{n1}^2 \ x_{n2}^2 \ x_{n1}x_{n2}]^T \in \mathbb{R}^6$
- The target t_n is regressed from the kernel representation $\phi(\mathbf{x}_n)$ as


$$\hat{t}_n = \mathbf{w}^T \phi(\mathbf{x}_n) \quad \mathbf{w} \in \mathbb{R}^{M+1}$$

- The regression coefficients are given by $\mathbf{w}_* = (\Phi^T \Phi)^{-1} \Phi^T \mathbf{t}$
 - $\Phi \in \mathbb{R}^{N \times M+1}$ denotes nonlinearly transformed data matrix
- DNNs can be used to learn data-dependent nonlinear transf. $\phi(\mathbf{x}_n)$
- The last layer of DNNs typically performs linear regression on $\phi(\mathbf{x}_n)$

Effect of Model Order M : $t = \sin(\pi x) + \epsilon$





Model Validation



	$M = 0$	$M = 1$	$M = 6$	$M = 9$
w_0^*	0.19	0.82	0.31	0.35
w_1^*		-1.27	7.99	232.37
w_2^*			-25.43	-5321.83
w_3^*			17.37	48568.31
w_4^*				-231639.30
w_5^*				640042.26
w_6^*				-1061800.52
w_7^*				1042400.18
w_8^*				-557682.99
w_9^*				125201.43

- Training & test error diverge for higher model orders
- Model 'overfits' to the noise in the training data
- Large amplitude weights with alternating polarity.
- $(\Phi^T \Phi)$ may be ill conditioned

Amount of Training Data ($M = 9$)



- Overfitting is less severe with increased amount of data.
- Model order cannot be limited by the amount of data available!
- Model order should be based on complexity of task/pattern!
- A way forward: arrest the growth of the model weights

Regularized Least Squares

- Add a penalty term to the error term to discourage weight growth

$$J(\mathbf{w}) = \underbrace{E_D(\mathbf{w})}_{\text{Data Term}} + \underbrace{\lambda E_W(\mathbf{w})}_{\text{Regularization Term}}$$

- λ controls the relative importance of the data and reg. terms
 - Smaller λ results in high variance, higher λ results in high bias
- Sum of squares error function with a quadratic regularizer

$$J(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N \left(t_n - \mathbf{w}^T \phi(\mathbf{x}_n) \right)^2 + \frac{\lambda}{2} \mathbf{w}^T \mathbf{w}$$


- Equating $\nabla_{\mathbf{w}} J(\mathbf{w}) = \mathbf{0} \implies -\Phi^T (\mathbf{t} - \Phi \mathbf{w}) + \lambda \mathbf{w} = \mathbf{0}$

$$\mathbf{w}_* = (\Phi^T \Phi + \lambda \mathbf{I})^{-1} \Phi^T \mathbf{t}$$


- Regularization term conditions the autocorrelation matrix!

Modified Error Surface

$$J(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N \left(t_n - \mathbf{w}^T \phi(\mathbf{x}_n) \right)^2 + \lambda \|\mathbf{w}\|_p$$

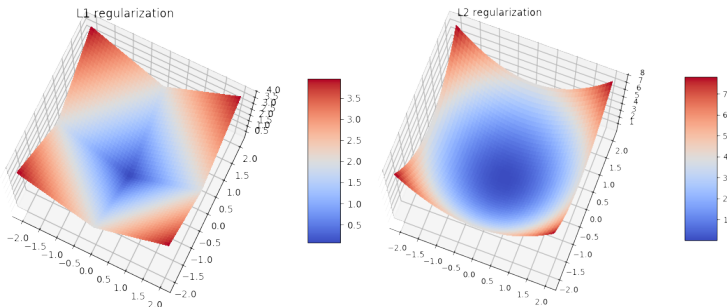


L_2 Regularizer





L_1 Regularizer

L_1 vs L_2




- L_2 regularization - also referred to as ridge regression
 - Promotes smaller and more evenly distributed weights
 - Equivalent to imposing Gaussian priors on parameters \mathbf{w}
- L_1 regularization, also referred to as LASSO
 - Promotes sparse models, which improves interpretability
 - Equivalent to imposing Laplacian priors on parameters \mathbf{w}


Effect of Regularization ($N = 10, M = 9$)



	$\ln \lambda = -\infty$	$\ln \lambda = -18$	$\ln \lambda = 0$
w_0^*	0.35	0.35	0.13
w_1^*	232.37	4.74	-0.05
w_2^*	-5321.83	-0.77	-0.06
w_3^*	48568.31	-31.97	-0.05
w_4^*	-231639.30	-3.89	-0.03
w_5^*	640042.26	55.28	-0.02
w_6^*	-1061800.52	41.32	-0.01
w_7^*	1042400.18	-45.95	-0.00
w_8^*	-557682.99	-91.53	0.00
w_9^*	125201.43	72.68	0.01



Issues with Least Squares Error



- Least squares error is highly sensitive to outliers
 - Squaring of errors amplifies the impact of large deviations
- LS assumes that errors are normally distributed with constant variance
- Computational inefficiency and potential numerical instability
 - For large datasets, pseudo-inverse can be computationally expensive
 - If features are highly correlated, then $\mathbf{X}^T \mathbf{X}$ becomes ill-conditioned

Regression - Loss Functions

- Least squares

$$J(\mathbf{w}) = \sum_{n=1}^N (t_n - y_n)^2$$

- Least absolute deviations

$$J(\mathbf{w}) = \sum_{n=1}^N |t_n - y_n|$$

- Huber loss

$$J(\mathbf{w}) = \sum_{n=1}^N \begin{cases} \frac{1}{2}(t_n - y_n)^2 & \text{if } |t_n - y_n| \leq \delta, \\ \delta|t_n - y_n| - \frac{1}{2}\delta^2 & \text{if } |t_n - y_n| > \delta. \end{cases}$$



- Log-Cosh loss

$$J(\mathbf{w}) = \sum_{n=1}^N \log(\cosh(t_n - y_n))$$

- KL Divergence

$$J(\mathbf{w}) = \sum_{n=1}^N t_n \log \frac{t_n}{y_n}$$

Illustration of Loss Functions




- LAD (L_1) is robust to outliers, but it is not differentiable around zero
- Log-Cosh and Huber follow L_2 at lower error and L_1 at higher errors
- For losses other than LS, we cannot get closed-form solutions

Sequential Learning

- The loss functions may not offer a closed-form solution
- In several applications, we may get data sequentially
- In such cases, we need to search for a solution on the error surface
- Iteratively update $\mathbf{w}^{(\tau+1)}$ by adding a correction factor

$$\mathbf{w}_{\tau+1} = \mathbf{w}_{\tau} + \Delta \mathbf{w}_{\tau}$$



Gradient Descent Algorithm

- Apply correction factor in the negative direction of gradient of $J(\mathbf{w}_\tau)$

$$\mathbf{w}_{\tau+1} = \mathbf{w}_\tau - \eta \nabla J(\mathbf{w})|_{\mathbf{w}=\mathbf{w}_\tau}$$

- η is the step-size parameter - controls the rate of convergence
- For least squares loss, the weight updates are given by

$$J(\mathbf{w}) = \frac{1}{2} \sum_n \left(t_n - \mathbf{w}^\top \phi(\mathbf{x}_n) \right)^2 \quad \nabla J(\mathbf{w}) = - \sum_n \left(t_n - \mathbf{w}^\top \phi(\mathbf{x}_n) \right) \phi(\mathbf{x}_n)$$

$$\mathbf{w}_{\tau+1} = \mathbf{w}_\tau + \eta \sum_n \left(t_n - \mathbf{w}_\tau^\top \phi(\mathbf{x}_n) \right) \phi(\mathbf{x}_n) \quad \tau = 0, 1, \dots$$

- For least absolute deviations, the weight updates are

$$J(\mathbf{w}) = \sum_n \left| t_n - \mathbf{w}^\top \phi(\mathbf{x}_n) \right| \quad \nabla J(\mathbf{w}) = - \sum_n \text{sign} \left(t_n - \mathbf{w}^\top \phi(\mathbf{x}_n) \right) \phi(\mathbf{x}_n)$$

$$\mathbf{w}_{\tau+1} = \mathbf{w}_\tau + \eta \sum_n \text{sign} \left(t_n - \mathbf{w}_\tau^\top \phi(\mathbf{x}_n) \right) \phi(\mathbf{x}_n)$$

Effect of Step-Size η


$$J(w) = (w-5)^2 + 20 \quad \nabla J(w) = 2(w-5)$$

$$w_{\tau+1} = w_{\tau} - 2\eta(w_{\tau} - 5)$$


$$w_0 = 25 \quad \& \quad \eta = 0.2$$

Evaluate w_{τ} for $\tau = 1, 2, 3$


$$w_0 = 25 \quad \eta = 0.2$$



$$\eta = 0.9$$



$$\eta = 0.05$$



Steepest Descent

- Full dataset is used to compute the gradient of the loss function
- The *true gradient* for the full data set is computed as

$$\nabla J(\mathbf{w}) = \sum_{n=1}^N \nabla J_n(\mathbf{w})$$

where $J_n(w)$ denotes loss on the n^{th} data point

- Error surface is consistent across the iterations
- Steepest descent converges smoothly for convex error surfaces

$$\lim_{\tau \rightarrow \infty} \mathbf{w}_\tau \rightarrow (\Phi^T \Phi)^{-1} \Phi^T \mathbf{t} \quad \text{for LS}$$

- Steepest descent gets stuck in local minima on nonconvex surfaces
- Gradient computation for all data at each iteration is expensive

Stochastic Gradient Descent (SGD)

- A random batch of points is used to compute the gradient
- Error surface slightly differs across iterations depending on batch
- *Noisy gradient* for a batch of points (\mathcal{B}) is computed as

$$\nabla J(\mathbf{w}) = \sum_{n \in \mathcal{B}} \nabla J_n(\mathbf{w})$$





where $|\mathcal{B}|$ is referred to as size of the batch. ($|\mathcal{B}| = 1$: LMS)

- The noisy gradients can help escape local minima and saddle points
- Suitable for very large datasets as full data is not processed at once
- Loss reduces quickly in the initial iterations but oscillates near minima

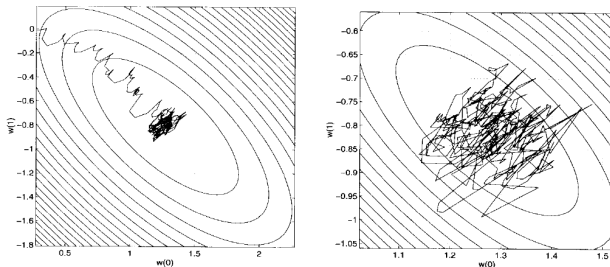
SGD Error Dynamics

$$t_n = w_1 x_{n1} + w_2 x_{n2} + \epsilon$$

$$w_1 = 1.6, w_2 = -0.5$$



Convergence of SGD







- SGD algorithm converges in mean:

$$\lim_{\tau \rightarrow \infty} \mathbb{E}[\mathbf{w}_\tau] \rightarrow (\Phi^T \Phi)^{-1} \Phi^T \mathbf{t} \quad \eta \text{ is small enough}$$

- Expectation over multiple runs (k) converges to true solution for convex error surfaces, provided η is sufficiently small

Geometry of Error Surface vs Convergence Rate



- Gradient magnitude depends on direction!
- η has to be fixed based on steepest direction.
- Convergence along flatter dimension is too slow!
- Normalization

$$\hat{\phi}_j(\mathbf{x}) = \frac{\phi_j(\mathbf{x}) - \mu_j}{\sigma_j}$$

Newton's Method

- The weights of the model are updated as

$$\mathbf{w}_{n+1} = \mathbf{w}_n + \Delta \mathbf{w}$$

- Expanding the objective function using Taylor series

$$J(\mathbf{w}_{n+1}) = J(\mathbf{w}_n + \Delta \mathbf{w}) = J(\mathbf{w}_n) + \Delta \mathbf{w}^T \nabla J(\mathbf{w}_n) + \frac{1}{2} \Delta \mathbf{w}^T \nabla^2 J(\mathbf{w}_n) \Delta \mathbf{w}$$

- Estimate $\Delta \mathbf{w}$ s.t $J(\mathbf{w}_n + \Delta \mathbf{w})$ is minimized

$$\frac{\partial}{\partial \Delta \mathbf{w}} \left(J(\mathbf{w}_n) + \Delta \mathbf{w}^T \nabla J(\mathbf{w}_n) + \frac{1}{2} \Delta \mathbf{w}^T \nabla^2 J(\mathbf{w}_n) \Delta \mathbf{w} \right) = 0$$

- Optimal update is given by $\Delta \mathbf{w} = -\frac{\nabla J(\mathbf{w}_n)}{\nabla^2 J(\mathbf{w}_n)}$

$$\mathbf{w}_{n+1} = \mathbf{w}_n - \mathbf{H}^{-1}(\mathbf{w}_n) \nabla J(\mathbf{w}_n) \quad \mathbf{H}(\mathbf{w}_n) = \nabla^2 J(\mathbf{w}_n)$$

Homework - 1

- Apply Newtons method to steepest-descent algorithm to the optimal step size η , and check how many iterations are required for convergence.

$$\mathbf{w}^{new} = \mathbf{w}^{old} + \eta \mathbf{X}^T (\mathbf{t} - \mathbf{X}\mathbf{w}) \Big|_{\mathbf{w}=\mathbf{w}^{old}}$$

Homework - 2

- Suppose you are experimenting with L_1 and L_2 regularization. Further, imagine that you are running gradient descent and at some iteration your weight vector is $w = [1, \epsilon] \in \mathbb{R}^2$ where $\epsilon > 0$ is very small. With the help of this example explain why L_2 norm does not encourage sparsity i.e., it will not try to drive ϵ to 0 to produce a sparse weight vector. Give mathematical explanation.

Homework - 3

- Till now we have been considering a scalar target t from a vector of input observations \mathbf{x} . How do you extend this approach for regressing a vector of targets $\mathbf{t} = (t_1, t_2, \dots, t_P)$. Derive the closed form solutions and write sequential update equations using SGD.

Probabilistic Approach to Regression

- Predict target variable(s) $t \in \mathbb{R}$ given the observation vector $\mathbf{x} \in \mathbb{R}^D$
- Target is estimated as a deterministic function $y(\mathbf{x}_n, \mathbf{w})$ with error e_n .


$$t_n = y(\mathbf{x}_n, \mathbf{w}) + e_n$$

- Assuming that the error is Gaussian distributed: $e_n \sim \mathcal{N}(0, \beta^{-1})$
 - The conditional distribution of target t_n is given by


$$p(t_n/\mathbf{x}_n, \mathbf{w}, \beta) \sim \mathcal{N}(y(\mathbf{x}_n, \mathbf{w}), \beta^{-1})$$

- β is referred to as the precision parameter
- We need to estimate \mathbf{w} and β to maximize $p(t_n/\mathbf{x}_n, \mathbf{w}, \beta) \quad \forall n$

Conditional Density of Target



Error Density



Target Conditional

- Gaussian error $e_n \implies$ Gaussian conditional density on targets t_n
- In this case, the mean of the target density is conditioned on input.
 - The observed samples are drawn from this conditional density
 - Given the observations, estimate the conditioning parameters

Maximum Likelihood (ML) Formulation


- Training data: $\{(\mathbf{x}_1, t_1), (\mathbf{x}_2, t_2) \cdots (\mathbf{x}_n, t_n) \cdots (\mathbf{x}_N, t_N)\}$
- Let the target be estimated as $\hat{t}_n = \mathbf{y}(\mathbf{x}_n, \mathbf{w}) = \mathbf{w}^T \phi(\mathbf{x}_n)$
- Assuming Gaussian errors: $p(t_n/\mathbf{x}_n, \mathbf{w}, \beta) = \mathcal{N}(\mathbf{w}^T \phi(\mathbf{x}_n), \beta^{-1})$
- Assuming the data points are drawn independently and identically

$$p(t_1, t_2, \cdots t_N/\mathbf{x}_1, \mathbf{x}_2, \cdots \mathbf{x}_N, \mathbf{w}, \beta) = \prod_{n=1}^N p(t_n/\mathbf{x}_n, \mathbf{w}, \beta)$$

$$\begin{aligned}\log p(\mathbf{t}/\mathbf{X}, \mathbf{w}, \beta) &= \sum_{n=1}^N \log \mathcal{N}(\mathbf{w}^T \phi(\mathbf{x}_n), \beta^{-1}) \\ &= \frac{N}{2} \log \beta - \frac{N}{2} \log(2\pi) - \frac{\beta}{2} \sum_{n=1}^N (t_n - \mathbf{w}^T \phi(\mathbf{x}_n))^2\end{aligned}$$


- \mathbf{w} and β have to be estimated to maximize likelihood $p(\mathbf{t}/\mathbf{X}, \mathbf{w}, \beta)$

Probability vs Likelihood



What is the probability that a sample drawn from a standard Gaussian lies between 1 & 2?

$$P[x_1 \leq X \leq x_2] = \int_{x_1}^{x_2} p_X(x) dx$$



Given that $x = 0.5$ is observed, is it likely to come from $p_{X_1}(x) = \mathcal{N}(0, 1)$ or $p_{X_2}(x) = \mathcal{N}(10, 1)$?


The value of the density function itself is considered the likelihood.

$$\mathcal{L}(\theta_1/x = 0.5) = p(x = 0.5/\theta_1)$$

Understanding Likelihood

- **Probability** of an event happening given the model (parameters)
 - Given the **model**, guess outcome
- **Likelihood** of a model (parameters) given the data
 - Given the **data**, assess how likely different models are.
- Given data D and parameters of a model θ , likelihood of model is

$$\mathcal{L}(\theta/D) = p(D/\theta)$$



ML \iff Least Squares

- ML with Gaussian conditional density assumption is same as LS

$$\mathbf{w}_{ML} = (\Phi^T \Phi)^{-1} \Phi^T \mathbf{t} \quad \frac{1}{\beta_{ML}} = \frac{1}{N} \sum_{n=1}^N \left(t_n - \mathbf{w}_{ML}^T \phi(\mathbf{x}_n) \right)^2$$

- ML approach assigns a probability density to the estimated target

$$p(t/\mathbf{x}, \mathbf{w}_{ML}, \beta_{ML}) = \mathcal{N}(y(\mathbf{x}, \mathbf{w}_{ML}), \beta_{ML}^{-1})$$

$$\mathbb{E}[t/\mathbf{x}] = \int t p(t/\mathbf{x}) dt = y(\mathbf{x}, \mathbf{w}_{ML})$$

- ML with Laplacian conditional density assumption is same as LAD
- ML & LS rely on point estimates of model parameters \mathbf{w}
- Point estimates cannot be exact with a finite number of samples
- Instead, estimate the distribution of \mathbf{w}

Maximum A Posteriori (MAP) Estimate

- Given a set of N datapoints, the posterior distribution of \mathbf{w} is

$$p(\mathbf{w}/\mathbf{t}, \mathbf{X}) \propto p(\mathbf{w})p(\mathbf{t}/\mathbf{w}, \mathbf{X})$$

- Let the prior distribution of \mathbf{w} be Gaussian: $p(\mathbf{w}) \sim \mathcal{N}(\mathbf{0}, \alpha^{-1}\mathbf{I})$
- Let the conditional distribution of target be Gaussian

$$p(t_n/\mathbf{x}_n, \mathbf{w}, \beta) \sim \mathcal{N}(\mathbf{w}^T \phi(\mathbf{x}_n), \beta^{-1})$$

- The posterior distribution of \mathbf{w} is given by

$$p(\mathbf{w}/\mathbf{t}, \mathbf{X}, \alpha, \beta) \propto \mathcal{N}(\mathbf{w}/\mathbf{0}, \alpha^{-1}\mathbf{I}) \prod_{n=1}^N \mathcal{N}(\mathbf{w}^T \phi(\mathbf{x}_n), \beta^{-1})$$

$$\log p(\mathbf{w}/\mathbf{t}) = -\frac{\beta}{2} \sum_{n=1}^N \left(t_n - \mathbf{w}^T \phi(\mathbf{x}_n) \right)^2 - \frac{\alpha}{2} \mathbf{w}^T \mathbf{w} + \text{const}$$

- Estimate \mathbf{w} to maximize $\log p(\mathbf{w}/\mathbf{t})$

MAP \iff Regularized Least Squares

- MAP estimation is equivalent to RLS with $\lambda = \frac{\alpha}{\beta}$
- MAP estimate of \mathbf{w} is given by

$$\mathbf{w}_{MAP} = (\Phi^T \Phi + \lambda \mathbf{I})^{-1} \Phi^T \mathbf{t}$$

- Gaussian priors $\iff L_2$ regularizer
- Laplacian priors $\iff L_1$ regularizer

Evaluating Posterior Density $p(\mathbf{w}/\mathbf{t})$

- Let the prior distribution of \mathbf{w} be $p(\mathbf{w}) = \mathcal{N}(\mathbf{w}/\mathbf{m}_0, \mathbf{\Sigma}_0)$
- Assuming linear model with Gaussian errors, the likelihood is given by

$$p(\mathbf{t}/\mathbf{X}, \mathbf{w}, \beta) = \prod_{n=1}^N \mathcal{N}(t_n/\mathbf{w}^T \phi(\mathbf{x}_n), \beta^{-1}) = \mathcal{N}(\mathbf{t}/\mathbf{\Phi w}, \beta^{-1} \mathbf{I})$$

- The posterior density after observing 'N' samples is given by

$$p(\mathbf{w}/\mathbf{t}) \propto \mathcal{N}(\mathbf{w}/\mathbf{m}_0, \mathbf{\Sigma}_0) \mathcal{N}(\mathbf{t}/\mathbf{\Phi w}, \beta^{-1} \mathbf{I}) = \mathcal{N}(\mathbf{w}/\mathbf{m}_N, \mathbf{\Sigma}_N)$$

- \mathbf{m}_N and $\mathbf{\Sigma}_N$ can be evaluated by completing quadratic term of $\exp()$

$$\mathbf{m}_N = \mathbf{\Sigma}_N \left(\mathbf{\Sigma}_0^{-1} \mathbf{m}_0 + \beta \mathbf{\Phi}^T \mathbf{t} \right)$$

$$\mathbf{\Sigma}_N^{-1} = \mathbf{\Sigma}_0^{-1} + \beta \mathbf{\Phi}^T \mathbf{\Phi}$$

Bayesian Sequential Estimates

- Let the posterior distribution of \mathbf{w} after observing n samples be

$$p(\mathbf{w}/\mathbf{t}_{1:n}) \sim \mathcal{N}(\mathbf{w}/\mathbf{m}_n, \mathbf{\Sigma}_n)$$

- In sequential update, $p(\mathbf{w}/\mathbf{t}_{1:n})$ is used as prior for $(n+1)^{th}$ sample
- The posterior stats can be updated after observing $(\mathbf{x}_{n+1}, t_{n+1})$ as

$$\mathbf{m}_{n+1} = \mathbf{\Sigma}_{n+1} (\mathbf{\Sigma}_n^{-1} \mathbf{m}_n + \beta \phi(\mathbf{x}_{n+1}) t_{n+1})$$

$$\mathbf{\Sigma}_{n+1}^{-1} = \mathbf{\Sigma}_n^{-1} + \beta \phi(\mathbf{x}_{n+1}) \phi^T(\mathbf{x}_{n+1})$$

Bayes Updates Illustration: $t = a_0 + a_1x + \epsilon$

- Actual targets are generated as


$$t = 0.5x - 0.3 + \epsilon$$

$$x \in \mathcal{U}[-1 \ 1] \quad \epsilon \in \mathcal{N}(0, 0.2^2)$$

- Assume: $y(x, \mathbf{w}) = w_1x + w_0$
- Assume noise variance is known

$$\beta = \frac{1}{0.2^2} \quad \alpha = 2.0$$

- Seq. update posterior $p(\mathbf{w}/\mathbf{t})$
- Draw random samples from $p(\mathbf{w}/\mathbf{t})$ and plot $y = w_1x + w_0$
- Lines converge as data increase



Homework

- Given a Gaussian marginal distribution for \mathbf{x} and a Gaussian conditional distribution for \mathbf{y} in the form

$$p(\mathbf{x}) = \mathcal{N}(\mathbf{x}/\boldsymbol{\mu}, \boldsymbol{\Lambda}^{-1})$$

$$p(\mathbf{y}/\mathbf{x}) = \mathcal{N}(\mathbf{y}/\mathbf{Ax} + \mathbf{b}, \mathbf{L}^{-1})$$

show that the marginal distribution of \mathbf{y} and conditional distribution of \mathbf{x} are given by

$$p(\mathbf{y}) = \mathcal{N}(\mathbf{y}/\mathbf{A}\boldsymbol{\mu} + \mathbf{b}, \mathbf{L}^{-1} + \mathbf{A}\boldsymbol{\Lambda}^{-1}\mathbf{A}^T)$$

$$p(\mathbf{x}/\mathbf{y}) = \mathcal{N}(\mathbf{x}/\boldsymbol{\Sigma}(\mathbf{A}^T\mathbf{L}(\mathbf{y} - \mathbf{b}) + \boldsymbol{\Lambda}\boldsymbol{\mu}), \boldsymbol{\Sigma})$$

$$\text{where } \boldsymbol{\Sigma} = (\boldsymbol{\Lambda} + \mathbf{A}^T\mathbf{L}\mathbf{A})^{-1}$$

Predictive Distributions

- Given a training set of N points $(\mathbf{x}_{1:N}, t_{1:N})$, predict target distribution for a new input \mathbf{x}_0

$$\begin{aligned} p(t_0/\mathbf{x}_0, \mathbf{X}, \mathbf{t}, \alpha, \beta) &= \int p(t_0, \mathbf{w}/\mathbf{x}_0, \mathbf{X}, \mathbf{t}, \alpha, \beta) d\mathbf{w} \\ &= \int p(t_0/\mathbf{w}, \mathbf{x}_0, \beta) p(\mathbf{w}/\mathbf{X}, \mathbf{t}, \alpha, \beta) d\mathbf{w} \end{aligned}$$

- The predictive distribution is Gaussian and is given by

$$\begin{aligned} p(t_0/\mathbf{x}_0, \mathbf{X}, \mathbf{t}, \alpha, \beta) &= \mathcal{N}\left(t_0/\mathbf{m}_N^T \phi(\mathbf{x}_0), \sigma_N^2(\mathbf{x}_0)\right) \\ \sigma_N^2(\mathbf{x}_0) &= \frac{1}{\beta} + \phi^T(\mathbf{x}_0) \mathbf{\Sigma}_N \phi(\mathbf{x}_0) \end{aligned}$$

- What happens to predictive distribution if we have one more data point?

Relation Between σ_{N+1}^2 and σ_N^2

- The predictive distribution after observing $N + 1$ points is

$$\sigma_{N+1}^2(\mathbf{x}_0) = \frac{1}{\beta} + \phi^\top(\mathbf{x}_0)\mathbf{\Sigma}_{N+1}\phi(\mathbf{x}_0)$$

- From Bayesian sequential estimates, $\mathbf{\Sigma}_{N+1}$ is related to $\mathbf{\Sigma}_N$ as

$$\mathbf{\Sigma}_{N+1}^{-1} = \mathbf{\Sigma}_N^{-1} + \beta\phi(\mathbf{x}_{N+1})\phi^\top(\mathbf{x}_{N+1})$$

- Woodbury matrix inversion lemma





$$(\mathbf{M} + \mathbf{v}\mathbf{v}^\top)^{-1} = \mathbf{M}^{-1} - \frac{\mathbf{M}^{-1}\mathbf{v}\mathbf{v}^\top\mathbf{M}^{-1}}{1 + \mathbf{v}^\top\mathbf{M}^{-1}\mathbf{v}}$$

- Predictive distribution gets narrower with additional training points





$$\sigma_{N+1}^2(\mathbf{x}_0) \leq \sigma_N^2(\mathbf{x}_0) \quad \lim_{N \rightarrow \infty} \sigma_N^2(\mathbf{x}_0) \rightarrow \frac{1}{\beta}$$

- Predictive uncertainty reduces with more data

Predictive Distribution: $t = \sin(2\pi x) + \epsilon$



Curves $y(x, \mathbf{w})$ Sampled from Posterior $p(\mathbf{w}/\mathbf{t})$



Summary of Linear Models of Regression

- Linear in model parameters \mathbf{w} .
 - If \mathbf{x} and t are linearly related $\hat{t} = \mathbf{w}^T \mathbf{x}$
 - If relationship is not linear: $\hat{t} = \mathbf{w}^T \phi(\mathbf{x})$
 - LS criterion leads to pseudo-inverse solution: $\mathbf{w}_* = (\Phi^T \Phi)^{-1} \Phi^T \mathbf{t}$
 - Regularize \mathbf{w} to avoid over-fitting: $\mathbf{w}_* = (\Phi^T \Phi + \lambda \mathbf{I})^{-1} \Phi^T \mathbf{t}$
 - Gradient descent algorithms can be used for sequential learning
- Probabilistic interpretation to regression
 - Point estimate of target does not hold for one-to-many maps
 - ML estimation assigns a distribution to the target $p(t_n/y(\mathbf{w}, \mathbf{x}_n), \beta^{-1})$
 - Parameters \mathbf{w} depend on training set - point estimate not enough
 - MAP estimation assigns a distribution to \mathbf{w} : $p(\mathbf{w}/\mathbf{t}, \mathbf{X}, \alpha, \beta)$
 - Predict target distribution for a test-point x_0 : $p(t_0/x_0, \mathbf{t}, \mathbf{X}, \alpha, \beta)$
 - Predictive uncertainty depends on x_0 and is smallest in the neighborhood of train data points.

Homework

- For Gaussian likelihood and Gaussian posterior, prove that the predictive distribution is Gaussian and is given by

$$p(t_0/\mathbf{x}_0, \mathbf{X}, \mathbf{t}, \alpha, \beta) = \mathcal{N}\left(t_0/\mathbf{m}_N^T \phi(\mathbf{x}_0), \sigma_N^2(\mathbf{x}_0)\right)$$

$$\sigma_N^2(\mathbf{x}_0) = \frac{1}{\beta} + \phi^T(\mathbf{x}_0) \mathbf{\Sigma}_N \phi(\mathbf{x}_0)$$

- Prove that the predictive uncertainty decreases with increase in training data, i.e., predictive distribution gets narrower with additional training points

$$\sigma_{N+1}^2(\mathbf{x}_0) \leq \sigma_N^2(\mathbf{x}_0)$$

Thank You!